

Assignment - 2
Data Visualization and
pre-processing

Assignment submission	26 September 2022
Student Name	Prasanth S
Student Roll Number	951919CS068
Maximum Marks	2 Marks

1. Download the Dataset

2. Import required library

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
sns.set_style('darkgrid')
sns.set(font_scale=1.3)
```

3. Read dataset and do pre-processing

```
df=pd.read_csv("/content/drive/MyDrive/IBM/Assignment - 2 /Churn_Modelling.csv")
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

Drop the columns that are not required for the neural network.


```
df.drop(["RowNumber","CustomerId","Surname"],axis=1,inplace=True)
```

```
df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
#   Column             Non-Null Count  Dtype  
---  -
0   CreditScore         10000 non-null  int64  
1   Geography           10000 non-null  object  
2   Gender              10000 non-null  object  
3   Age                 10000 non-null  int64  
4   Tenure              10000 non-null  int64  
5   Balance             10000 non-null  float64 
6   NumOfProducts       10000 non-null  int64  
7   HasCrCard           10000 non-null  int64  
8   IsActiveMember      10000 non-null  int64  
9   EstimatedSalary     10000 non-null  float64 
10  Exited              10000 non-null  int64  
dtypes: float64(2), int64(7), object(2)
memory usage: 859.5+ KB
```

4.A. Perform Univariate Analysis

```
plt.figure(figsize=(8,8))
sns.countplot(x='Tenure',data=df)
plt.xlabel('0:Customers with Bank, 1: exited from bank')
plt.ylabel('No.of.Customers')
plt.title("Bank Customers viz")
plt.show()
```




```
plt.figure(figsize=(8,8))
sns.kdeplot(x=df['Balance'])
```



4.B. Perform Bi-variate Analysis

```
plt.scatter(df.Age,df.Balance)
```

<matplotlib.collections.PathCollection at 0x7fa0d35a7dd0>



```
df.corr()
```

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
CreditScore	1.000000	0.007888	-0.003965	0.000842	0.006268	0.012238	-0.005458	0.025651	-0.001384	-0.027094
Gender	0.007888	1.000000	0.022812	0.003739	0.069408	0.003972	-0.008523	0.006724	-0.001369	0.035943
Age	-0.003965	0.022812	1.000000	-0.009997	0.028308	-0.030680	-0.011721	0.085472	-0.007201	0.285323
Tenure	0.000842	0.003739	-0.009997	1.000000	-0.012254	0.013444	0.022583	-0.028362	0.007784	-0.014001
Balance	0.006268	0.069408	0.028308	-0.012254	1.000000	-0.304180	-0.014858	-0.010084	0.012797	0.118533
NumOfProducts	0.012238	0.003972	-0.030680	0.013444	-0.304180	1.000000	0.003183	0.009612	0.014204	-0.047820
HasCrCard	-0.005458	-0.008523	-0.011721	0.022583	-0.014858	0.003183	1.000000	-0.011866	-0.009933	-0.007138
IsActiveMember	0.025651	0.006724	0.085472	-0.028362	-0.010084	0.009612	-0.011866	1.000000	-0.011421	-0.156128
EstimatedSalary	-0.001384	-0.001369	-0.007201	0.007784	0.012797	0.014204	-0.009933	-0.011421	1.000000	0.012097
Exited	-0.027094	0.035943	0.285323	-0.014001	0.118533	-0.047820	-0.007138	-0.156128	0.012097	1.000000

```
#Perform Bivariate Analysis
import statsmodels.api as sm
```

```
#define response variable
y = df['CreditScore']
```

```
#define explanatory variable
x = df[['EstimatedSalary']]
```

```
#add constant to predictor variables
x = sm.add_constant(x)
```

```
#fit linear regression model
model = sm.OLS(y, x).fit()
```

```
#view model summary
print(model.summary())
```

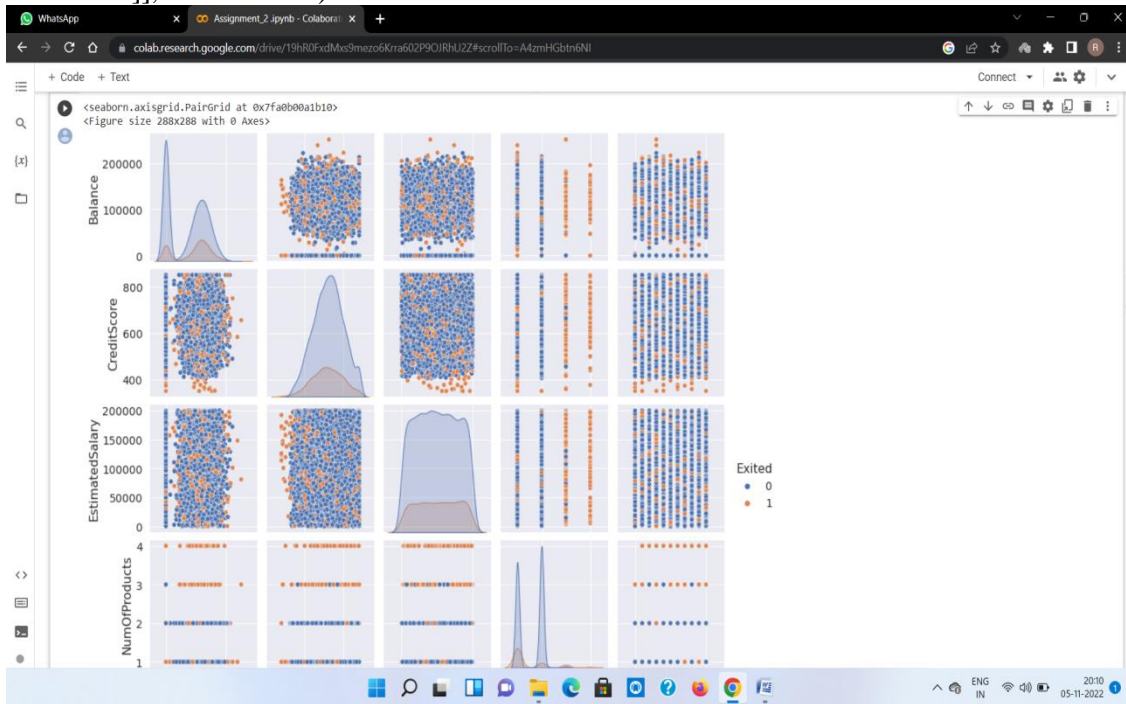
```
OLS Regression Results
=====
Dep. Variable:  CreditScore  R-squared:  0.000
Model:  OLS  Adj. R-squared:  -0.000
Method:  Least Squares  F-statistic:  0.01916
Date:  Sat, 24 Sep 2022  Prob (F-statistic):  0.890
Time:  05:06:19  Log-Likelihood:  -59900.
No. Observations:  10000  AIC:  1.198e+05
Df Residuals:  9998  BIC:  1.198e+05
Df Model:  1
Covariance Type:  nonrobust
=====
               coef      std err          t      P>|t|      [0.025   0.975]
-----
const      650.7617      1.940      335.407      0.000      646.958      654.565
EstimatedSalary -2.326e-06  1.68e-05      -0.138      0.890      -3.53e-05      3.06e-05
=====
Omnibus:  132.939  Durbin-Watson:  2.014
Prob(Omnibus):  0.000  Jarque-Bera (JB):  84.242
Skew:  -0.072  Prob(JB):  5.10e-19
Kurtosis:  2.574  Cond. No.  2.32e+05
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 2.32e+05. This might indicate that there are
strong multicollinearity or other numerical problems.
/usr/local/lib/python3.7/dist-packages/statsmodels/tsa/tsatools.py:142: FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be key
x = pd.concat(x[::order], 1)
```

4.C. Perform Multi-variate Analysis

#Perform Multivariate Analysis

```
plt.figure(figsize=(4,4))
sns.pairplot(data=df[["Balance", "CreditScore", "EstimatedSalary", "NumOfProducts", "Tenure",
"Exited"]], hue="Exited")
```



5.Perform descriptive statistics on the datasets:

```
df=pd.DataFrame(df)
print(df.sum())
```

```
CreditScore      6505288
Geography      FranceSpainFranceFranceSpainSpainFranceGermany...
Gender      FemaleFemaleFemaleFemaleFemaleMaleFemaleMaleFemale...
Age      389218
Tenure      50128
Balance      764858892.88
NumOfProducts      15302
HasCrcCard      7055
IsActiveMember      5151
EstimatedSalary      1000902398.81
Exited      2037
dtype: object
```

#Perform Descriptive Statistics

```
print("----Sum Value----")
print(df.sum(1))
print("-----")
print("----Product Value----")
print(df.prod())
print("-----")
```

```
[x]
-----Sum Value-----
0      102015.88
1      197002.44
2      274149.37
3      94567.63
4      205492.92
...
9995    97088.64
9996   159633.38
9997    42840.58
9998   168784.83
9999   169159.57
length: 10000, dtype: float64

-----Product Value-----
CreditScore    0.0
Age            0.0
Tenure         0.0
Balance        0.0
NumOfProducts 0.0
HasCrCard      0.0
IsActiveMember 0.0
EstimatedSalary inf
Exited         0.0
dtype: float64

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:3: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a futur
This is separate from the ipykernel package so we can avoid doing imports until
/usr/local/lib/python3.7/dist-packages/numpy/core/_methods.py:52: RuntimeWarning: overflow encountered in reduce
return umr.prod(a, axis, dtype, out, keepdims, initial, where)
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:6: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a futur
```

#Perform Descriptive Statistics

```
print("-----Mean Value-----")
print(df.mean())
print("-----")
print("-----Median Value-----")
print(df.median())
print("-----")
print("-----Mode Value-----")
print(df.mode())
print("-----")
```

```

-----Mean Value-----
CreditScore    650.528800
Age            38.921800
Tenure         5.012800
Balance        76485.889288
NumOfProducts  1.530200
HasCrCard      0.705500
IsActiveMember 0.515100
EstimatedSalary 100090.239881
Exited         0.203700
dtype: float64

-----Median Value-----
CreditScore    652.000
Age            37.000
Tenure         5.000
Balance        97198.540
NumOfProducts  1.000
HasCrCard      1.000
IsActiveMember 1.000
EstimatedSalary 100193.915
Exited         0.000
dtype: float64

-----Mode Value-----
CreditScore Geography Gender Age Tenure Balance NumOfProducts \
0      850      France   Male   37      2      0.0      1

HasCrCard  IsActiveMember  EstimatedSalary  Exited
0          1              1      24924.92      0

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:3: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a futur
This is separate from the ipykernel package so we can avoid doing imports until
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:6: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a futur
```

6.Handle the missing values:

df.isnull()#Checking values are null

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	False	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False
...
9995	False	False	False	False	False	False	False	False	False	False	False
9996	False	False	False	False	False	False	False	False	False	False	False
9997	False	False	False	False	False	False	False	False	False	False	False
9998	False	False	False	False	False	False	False	False	False	False	False
9999	False	False	False	False	False	False	False	False	False	False	False

10000 rows x 11 columns

#Handling with missing Values

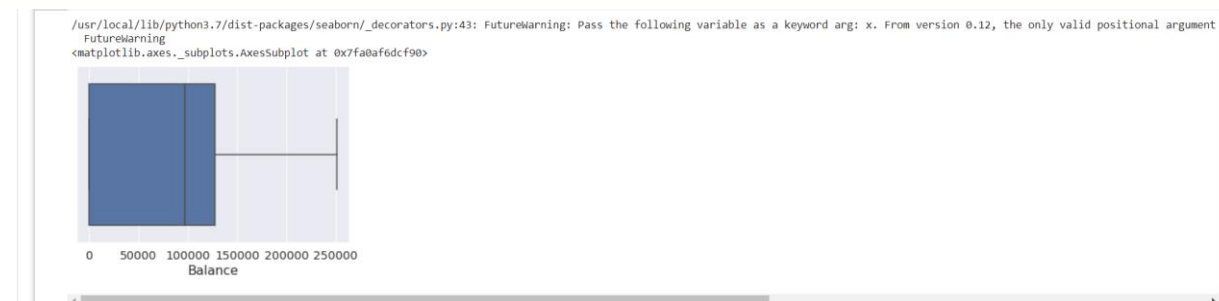
df.notnull()#Checking values are not null

	Creditscore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	True	True	True	True	True	True	True	True	True	True	True
1	True	True	True	True	True	True	True	True	True	True	True
2	True	True	True	True	True	True	True	True	True	True	True
3	True	True	True	True	True	True	True	True	True	True	True
4	True	True	True	True	True	True	True	True	True	True	True
...
9995	True	True	True	True	True	True	True	True	True	True	True
9996	True	True	True	True	True	True	True	True	True	True	True
9997	True	True	True	True	True	True	True	True	True	True	True
9998	True	True	True	True	True	True	True	True	True	True	True
9999	True	True	True	True	True	True	True	True	True	True	True

10000 rows x 11 columns

7.Find outlier and replace the outlier:

sns.boxplot(df['Balance'])



print(np.where(df['Balance']>100000))

(array([2, 4, 5, ..., 9987, 9993, 9999]),)

#Find outliers & replace the outliers

from scipy import stats

import numpy as np

z = np.abs(stats.zscore(df["EstimatedSalary"]))

print(z)

```
0    0.021886
1    0.216534
2    0.240687
3    0.108918
4    0.365276
...
9995  0.066419
9996  0.027988
9997  1.008643
9998  0.125231
9999  1.076370
Name: EstimatedSalary, Length: 10000, dtype: float64
```

8.Check for categorical columns & performs encoding:

from sklearn.preprocessing import LabelEncoder

df['Gender'].unique()

array(['Female', 'Male'], dtype=object)

df['Gender'].value_counts()

```
Male    5457
Female  4543
Name: Gender, dtype: int64
```

```
#Check for categorical columns & performs encoding
encoding=LabelEncoder()
df["Gender"]=encoding.fit_transform(df.iloc[:,1].values)
df
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	619	France	0	42	2	0.00	1	1	1	101348.88	1
1	608	Spain	2	41	1	83807.86	1	0	1	112542.58	0
2	502	France	0	42	8	159660.80	3	1	0	113931.57	1
3	699	France	0	39	1	0.00	2	0	0	93826.63	0
4	850	Spain	2	43	2	125510.82	1	1	1	79084.10	0
...
9995	771	France	0	39	5	0.00	2	1	0	96270.64	0
9996	516	France	0	35	10	57369.61	1	1	1	101699.77	0
9997	709	France	0	36	7	0.00	1	0	1	42085.58	1
9998	772	Germany	1	42	3	75075.31	2	1	0	92888.52	1
9999	792	France	0	28	4	130142.79	1	1	0	38190.78	0

10000 rows x 11 columns

9.Split the data into Dependent & Independent Variables:

```
print("-----Dependent Variables-----")
X=df.iloc[:,1:4]
print(X)
print("-----")
print("-----Independent Variables-----")
Y=df.iloc[:,4]
print(Y)
print("-----")
```

-----Dependent Variables-----		
	Age	Balance
0	42	2
1	41	1
2	42	8
3	39	1
4	43	2
...
9995	39	5
9996	35	10
9997	36	7
9998	42	3
9999	28	4

[10000 rows x 3 columns]

-----Independent Variables-----		
	NumOfProducts	HasCrCard
0	1	1
1	1	0
2	3	1
3	2	0
4	1	1
...
9995	2	1
9996	1	1
9997	1	0
9998	2	1
9999	1	1

Name: NumOfProducts, Length: 10000, dtype: int64

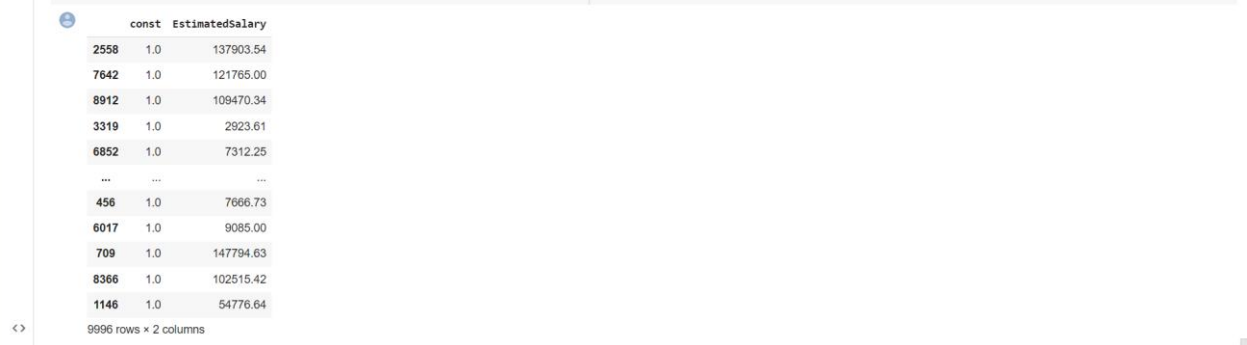
10.Scale the independent variables:

```
from sklearn.preprocessing import StandardScaler
object= StandardScaler()
# standardization
scale = object.fit_transform(df)
print(scale)
```

[[-0.32622142 0.29351742 -1.04175968 ... 0.97024255 0.02188649
1.97716468]
[-0.44003595 0.19816383 -1.38753759 ... 0.97024255 0.21653375
-0.50577476]
[-1.53679418 0.29351742 1.03290776 ... -1.03067011 0.2406869
1.97716468]
...
[0.60498839 -0.27860412 0.68712986 ... 0.97024255 -1.00864308
1.97716468]
[1.25683526 0.29351742 -0.69598177 ... -1.03067011 -0.12523071
1.97716468]
[1.46377078 -1.04143285 -0.35020386 ... -1.03067011 -1.07636976
-0.50577476]]

11. Split the data into training & testing:

from sklearn.model_selection import train_test_split

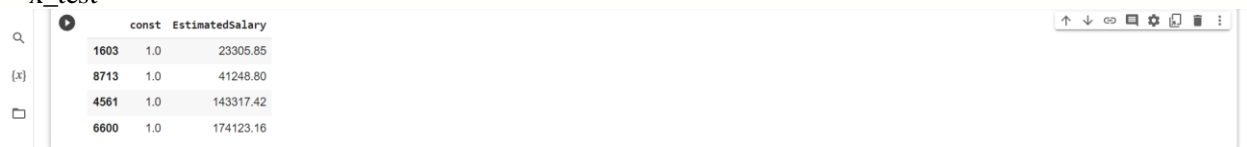


A Jupyter Notebook cell displaying a DataFrame with two columns: 'const' and 'EstimatedSalary'. The 'const' column contains values like 2558, 7642, 8912, 3319, 6852, ..., 456, 6017, 709, 8366, 1146. The 'EstimatedSalary' column contains corresponding salary values like 137903.54, 121765.00, 109470.34, 2923.61, 7312.25, ..., 7666.73, 9085.00, 147794.63, 102515.42, 54776.64. The bottom of the cell indicates '9996 rows x 2 columns'.

	const	EstimatedSalary
	2558	137903.54
	7642	121765.00
	8912	109470.34
	3319	2923.61
	6852	7312.25

	456	7666.73
	6017	9085.00
	709	147794.63
	8366	102515.42
	1146	54776.64

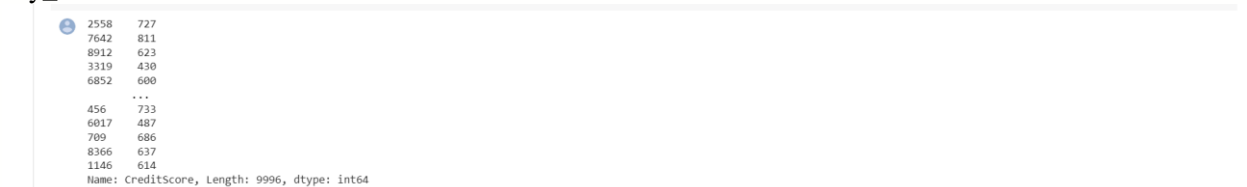
x_test



A Jupyter Notebook cell displaying a DataFrame with two columns: 'const' and 'EstimatedSalary'. The 'const' column contains values like 1603, 8713, 4561, 6600. The 'EstimatedSalary' column contains corresponding salary values like 23305.85, 41248.80, 143317.42, 174123.16.

	const	EstimatedSalary
	1603	23305.85
	8713	41248.80
	4561	143317.42
	6600	174123.16

y_train



A Jupyter Notebook cell displaying a Series of CreditScore values. The values are 727, 811, 623, 430, 600, ..., 733, 487, 686, 637, 614. The bottom of the cell indicates 'Name: CreditScore, Length: 9996, dtype: int64'.

	CreditScore
	727
	811
	623
	430
	600
	...
	733
	487
	686
	637
	614

y_test



A Jupyter Notebook cell displaying a Series of CreditScore values. The values are 576, 786, 562, 505. The bottom of the cell indicates 'Name: CreditScore, dtype: int64'.

	CreditScore
	576
	786
	562
	505